

12/09/2024



# Balancing Complexity and Accuracy: Predicting Stock Price Directions with Deep Learning

**Group 2**

**Deep Learning Final Project**

**Group Members**

Jianjun Gao

Luhuan Wang

# Data Source

## Financial Modeling Prep API

The screenshot shows the homepage of the Financial Modeling Prep (FMP) website. At the top, there is a navigation bar with the FMP logo, a search bar, and links for API Tools, Pricing, Blog, Screener, and Extras. A secondary bar displays various market indices and their percentage changes, such as DOW 43.12 ▲ 2.68%, BTCUSD 97633.97 ▼ -3.53564%, and AAPL 246.21 ▲ 1.39%.

The main content area features a large heading "Financial data for every need" with a subtext: "We provide one of the most accurate financial data available on the market. You can get historical prices, fundamental data, insider transactions, and much more that goes back 30 years in history." Below this is a "Get started" button.

To the right, a callout box displays a sample financial statement for Apple Inc. (AAPL) as of 9/30/17, showing a +0.06 (+1.84%) change. The statement includes metrics like Revenue (365,810,000), Net Income (94,680,000), and Earnings Per Share Basic (5.67).

Below the main text, it states "70,000+ Stocks with tons of data" and lists several companies: NETFLIX, twitter, Alphabet, amazon.com, and TESLA.

The footer section is titled "Ready to Unlock Insights?" and includes a welcome message: "Welcome to our comprehensive suite of tools designed to streamline your financial data integration process." It also features a "STOCK PRICES" section with a description: "Real-time Prices, Historical Data and Indicators." and a "MARKET NEWS" section.

# Data Source

7k Stocks from NASDAQ & NYSE(Excluding bonds, mutual funds, options, futures and etc.)

Changelog

Google Sheet Integration

Excel Add-on

Company Search

Stock List

Symbol List

Exchange Traded Fund Search

Statement Symbols List

Tradable Search

Commitment Of Traders Report

13F CIK List

Euronext Symbols

Symbol Changes

Exchange Symbols

Available Indexes

Company Information

Quote

Financial Statements

Income Statement

Balance Sheet Statement

Cashflow Statement

Income Statement As Reported

Balance Statement As Reported

Cashflow Statement As Reported

Full Financial Statement As Reported

List Of Dates

Annual Reports On Form 10-K

Statement Analysis

Valuation

Price Targets

Upgrades & Downgrades

Financial Statements

The financial statements, including balance sheet, income statement, and cash flow statement available in annual and quarterly format sourced from SEC filings

Income Statements API

FMP's Income Statement API provides access to real-time income statement data for a wide range of companies, including public companies, private companies, and ETFs. This data can be used to track a company's profitability over time, to compare a company to its competitors, and to identify trends in a company's business.

Endpoint:

https://financialmodelingprep.com/api/v3/income-statement/AAPL?period=annual

Endpoint:

https://financialmodelingprep.com/api/v3/income-statement/0000320193?period=annual

Parameters

Path Parameter	Type	Example
symbol *	string	AAPL
cik *	string	0000320193

Query Parameter	Type	Example
period	string	annual, quarter
datatype	string	csv
limit	number	100

Response

```
1 [
2   {
3     "date": "2022-09-24",
4     "symbol": "AAPL",
5     "reportedCurrency": "USD",
6     "cik": "0000320193",
7     "fillingDate": "2022-10-28",
8     "acceptedDate": "2022-10-27 18:01:14",
9     "calendarYear": "2022",
10    "period": "FY",
11    "revenue": 394328000000,
12    "costOfRevenue": 223546000000,
13    "grossProfit": 170782000000,
14    "grossProfitRatio": 0.4330963056,
15    "researchAndDevelopmentExpenses": 26251000000,
16    "generalAndAdministrativeExpenses": 0,
17    "sellingAndMarketingExpenses": 0,
18    "sellingGeneralAndAdministrativeExpenses": 25094000000,
```

# Data Cleaning

- Removing duplicates and NaN values.
- Removing all irrelevant features including 'Currency', 'cik', 'link', etc.
- Keep all the companies that have annual financial statements in recent 5 years.

# Reasons for choosing a 5-year time series length as model input


Shorter time series, such as 1 year, may ignore medium-term market trends and focus only on short-term fluctuations.

Longer time series, such as 10 years, may introduce too much redundant information and increase the complexity and computational cost of the model.

Five years of historical financial data and stock price information can usually reflect a company's medium-term trend, especially for annual report data. This length can help capture some characteristics of financial cycles. In data analysis and model experiments, an input length of 5 years may show the best performance on the validation set and therefore become the final choice.

# Data Features

Correlation Matrix



'Revenue'  
'grossProfit'  
'operatingIncome'  
'netIncome'  
'cashAndCashEquivalents'  
'totalAssets'  
'totalLiabilities'  
'sections'  
'adjClose'

# Feature Engineering

```
df['eps_normalized'] = df['eps'] / df['adjClose']  
df['debt_to_equity'] = df['totalDebt'] / (df['totalEquity'] + 1e-6)  
df['market_cap'] = df['adjClose'] * df['weightedAverageShsOut']
```

# Logic for generating price direction binary classification labels

Use the adjusted closing price (adjClose):

The adjusted closing price reflects the impact of dividends and stock splits on the price and is a more accurate price data.

Calculate the trend in the next year:

Shift the adjClose value forward (using `.shift(-1)`), that is, compare the price of the current year with the price of the next year.

If the adjusted closing price of the next year is higher than the current year, the label is 1, indicating an increase.

If the adjusted closing price of the next year is lower than the current year, the label is 0, indicating a decrease.



# Model architecture

Selected models

**RNN:** A model built with simple RNN layers.

**LSTM:** Bidirectional LSTM is used to capture long-term dependencies in time series.

**CNN-LSTM:** Convolutional layers extract features and LSTM is used to process sequence data.

**LSTM-Attention:** An Attention layer is added to LSTM for weight distribution.

# Experimental settings

Data split

Training set: 80%, test set: 20%.

Training parameters

Loss function: binary cross entropy.

Optimizer: AdamW (learning rate = 0.001).

Callback function: learning rate adjustment (ReduceLROnPlateau) and early stopping (EarlyStopping).

Hyperparameters

Sequence length: 5.

Batch size: 32.

Training epochs: up to 50 epochs.

# RNN

Training RNN model...


Model: "sequential\_1"

Layer (type)	Output Shape	Param #
simple_rnn (SimpleRNN)	(None, 5, 128)	17,536
dropout_3 (Dropout)	(None, 5, 128)	0
simple_rnn_1 (SimpleRNN)	(None, 64)	12,352
dropout_4 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 64)	4,160
dropout_5 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 32)	2,080
dense_5 (Dense)	(None, 1)	33

Total params: 36,161 (141.25 KB)

Trainable params: 36,161 (141.25 KB)

Non-trainable params: 0 (0.00 B)



# LSTM

```
Training LSTM model...
```

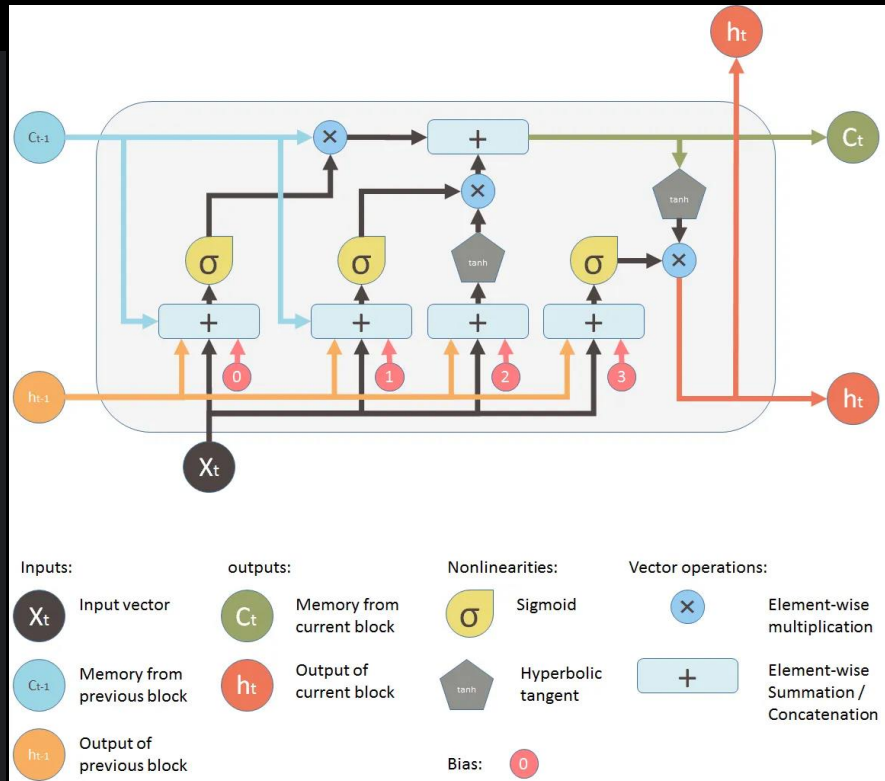
```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
bidirectional (Bidirectional)	(None, 5, 256)	140,288
dropout (Dropout)	(None, 5, 256)	0
lstm_1 (LSTM)	(None, 64)	82,176
dropout_1 (Dropout)	(None, 64)	0
dense (Dense)	(None, 64)	4,160
dropout_2 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2,080
dense_2 (Dense)	(None, 1)	33

Total params: 228,737 (893.50 KB)

Trainable params: 228,737 (893.50 KB)

```
Non-trainable params: 0 (0.00 B)
```



# CNN + LSTM

Training CNN-LSTM model...

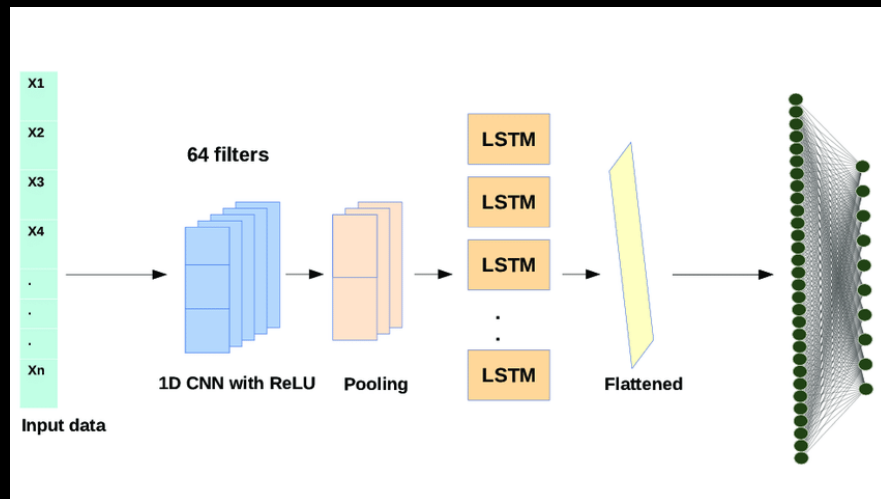
Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 3, 64)	1,600
max_pooling1d (MaxPooling1D)	(None, 1, 64)	0
dropout_6 (Dropout)	(None, 1, 64)	0
bidirectional_1 (Bidirectional)	(None, 128)	66,048
dropout_7 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 64)	8,256
dropout_8 (Dropout)	(None, 64)	0
dense_7 (Dense)	(None, 32)	2,080
dense_8 (Dense)	(None, 1)	33

Total params: 78,017 (304.75 KB)

Trainable params: 78,017 (304.75 KB)

Non-trainable params: 0 (0.00 B)



# LSTM-Attention

Training LSTM-Attention model...

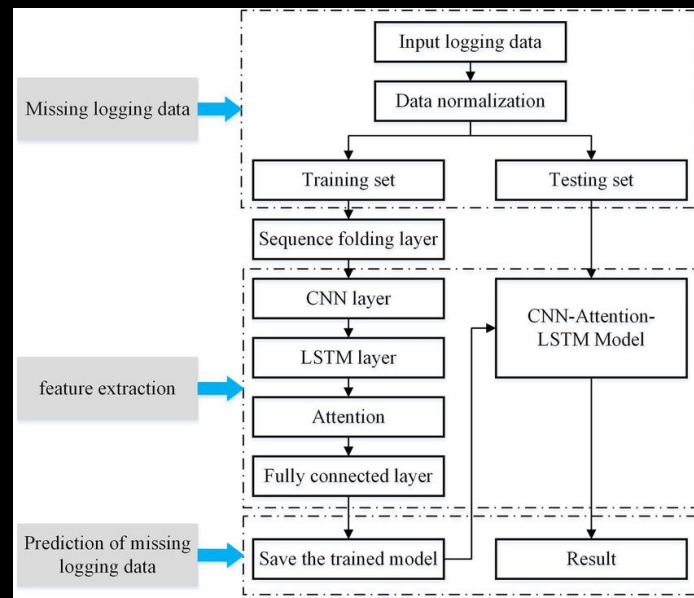
Model: "sequential\_3"

Layer (type)	Output Shape	Param #
bidirectional_2 (Bidirectional)	(None, 5, 256)	140,288
dropout_9 (Dropout)	(None, 5, 256)	0
lstm_4 (LSTM)	(None, 5, 64)	82,176
dropout_10 (Dropout)	(None, 5, 64)	0
attention (Attention)	(None, 64)	65
dense_9 (Dense)	(None, 64)	4,160
dropout_11 (Dropout)	(None, 64)	0
dense_10 (Dense)	(None, 32)	2,080
dense_11 (Dense)	(None, 1)	33

Total params: 228,802 (893.76 KB)

Trainable params: 228,802 (893.76 KB)

Non-trainable params: 0 (0.00 B)



```

class Attention(Layer):
    def __init__(self, **kwargs):
        super(Attention, self).__init__(**kwargs)

    def build(self, input_shape):
        self.W = self.add_weight(name='attention_weight', shape=(input_shape[-1], 1),
                                initializer='random_normal', trainable=True)
        self.b = self.add_weight(name='attention_bias', shape=(1,),
                                initializer='zeros', trainable=True)
        super(Attention, self).build(input_shape)

    def call(self, x):
        e = K.tanh(K.dot(x, self.W) + self.b)
        a = K.softmax(e, axis=1)
        output = x * a
        return K.sum(output, axis=1)
    
```

# Experimental results

Model performance comparison  
Verification accuracy and test set  
accuracy of each model.

A line chart is used to show the trend  
of verification accuracy of each  
model as the training round changes.


# Key points of the chart

RNN (orange line): The highest validation accuracy, stable performance, and small fluctuations.

LSTM (blue line): Performance is close to RNN, but with slight fluctuations.

CNN-LSTM (green): The validation accuracy is low and has never exceeded 65%.

LSTM-Attention (red line): It improves rapidly in the initial training, and then reach the highest accuracy of RNN.





# Final Test Accuracy Comparison

The RNN model performed best with a test accuracy of 66.50%, indicating that it has an advantage in capturing short- and medium-term features of time series.

The test accuracy of the LSTM model and LSTM-Attention model were 65.04% and 65.59%, respectively, which were similar but slightly lower than RNN.

The accuracy of the CNN-LSTM model was 64.71%, the worst among all models, indicating that its feature extraction ability was not ideal in this task.


Final Model Accuracies:

LSTM: 0.6504

RNN: 0.6605

CNN-LSTM: 0.6471

LSTM-Attention: 0.6559



# Comparative analysis of model performance

**Best model:** Based on experimental results, the LSTM-Attention model performs superiorly with its weight allocation mechanism and has significant advantages in capturing important time step features in time series.

Model performance comparison: The RNN model achieved the highest accuracy of 66.05% on the test set, but the LSTM-Attention model followed closely behind, with a more stable performance in the validation set and higher potential interpretability.

The CNN-LSTM model performs relatively poorly, possibly due to poor compatibility between convolutional feature extraction and sequence tasks.

The LSTM model performs close to RNN in the validation and test sets, but slightly worse than the latter.

# Overall rating of the model

The LSTM-Attention model has stronger feature extraction capabilities with the support of the weight allocation mechanism, and is a deep learning framework suitable for this task.

Due to its simple structure, the RNN model can better handle short- and medium-term changes in current features.

# Reference:

<https://blog.mlreview.com/understanding-lstm-and-its-diagrams-37e2f46f1714>

[https://www.researchgate.net/figure/Architecture-of-the-Hybrid-1D-CNN-LSTM-model-for-human-activity-recognition\\_fig4\\_343341551](https://www.researchgate.net/figure/Architecture-of-the-Hybrid-1D-CNN-LSTM-model-for-human-activity-recognition_fig4_343341551)

<https://medium.com/@poudelsushmita878/recurrent-neural-network-rnn-architecture-explained-1d69560541ef>

[https://www.researchgate.net/figure/Flow-chart-of-CNN-LSTM-Attention-model\\_fig3\\_363533496](https://www.researchgate.net/figure/Flow-chart-of-CNN-LSTM-Attention-model_fig3_363533496)

<https://site.financialmodelingprep.com/>

Python Libraries and Tools:

TensorFlow/Keras for building deep learning models.

Pandas, NumPy, and Matplotlib for data preprocessing and visualization.

Scikit-learn for data splitting and scaling.

# Thank you all for your time and attention!

In this presentation, we explored the application of deep learning models for predicting stock price directions and conducted a comparative analysis of various models. The results demonstrated the potential of advanced techniques in financial forecasting, while also highlighting areas for improvement and future research opportunities.



Thank you!